

Storyboard Generation via Diffusion Models

Sandro Mikautadze

Elio Samaha

École Polytechnique, Paris

Abstract

Storyboarding constitutes a fundamental component in visual narrative construction, traditionally necessitating substantial manual effort and artistic proficiency. This research presents a novel AI-driven methodology for automating storyboard generation through the utilization of text-to-image diffusion models. Our framework employs a hierarchical approach, wherein scripts and characters are processed via large language models (LLMs) to generate sequential keyframe descriptions, from the generation of the images occurs. We investigate three diffusion-based image generation methodologies: (i) Unique Prompting, (ii) Weighted Embeddings, and (iii) Multi-Prompt Classifier-Free Guidance. Experimental results conducted on script excerpts demonstrate the relative efficacy, and advantages of each proposed approach. We identify current technical challenges, and propose directions for future investigation to address these limitations. The code is available at <https://github.com/sandromikautadze/storyboard-diffusion>.

1 Introduction

Visual narratives in cinematic and graphic novel production critically depend on storyboards—sequential visual keyframes that function as essential production blueprints. Conventional storyboard methodologies impose significant resource constraints due to extensive manual illustration requirements and iterative revision processes, frequently creating production bottlenecks. Recent advancements in text-conditional generative models, particularly diffusion-based architectures, present promising avenues for automating this labor-intensive process. In this work, we present a framework for automated storyboard generation that integrates several components:

1. **Hierarchical Narrative Decomposition:** Leveraging large language models (LLMs) to parse scripts and character descriptions into semantically structured storyboard panel representations.
2. **Diffusion-Based Image Synthesis:** Employing latent diffusion models (specifically, Stable Diffusion v1.5) to generate high-fidelity visualizations from structured textual prompts.
3. **Advanced Prompting Techniques:** Implementing weighted embedding compositions and multi-prompt classifier-free guidance to overcome text-length constraints while enhancing narrative and style.

Our results show a promising direction for future work. Overall, Section 3 does an overview of the methods applied for this work. We also show some other approaches we tried during the whole project in Section 4. Section 5 outlines the experiments along with their results. Finally in Sections 6 we conclude with a discussion on the outcomes of the experiments and on future improvements.

2 Background

- **Text-to-Image Synthesis:** One of the latest and biggest revolutions in text-to-image generation comes with **Stable Diffusion**[4], which demonstrates superior sample quality and training stability for high-resolution image generation compared to previous methods, such as GANs[1]. The

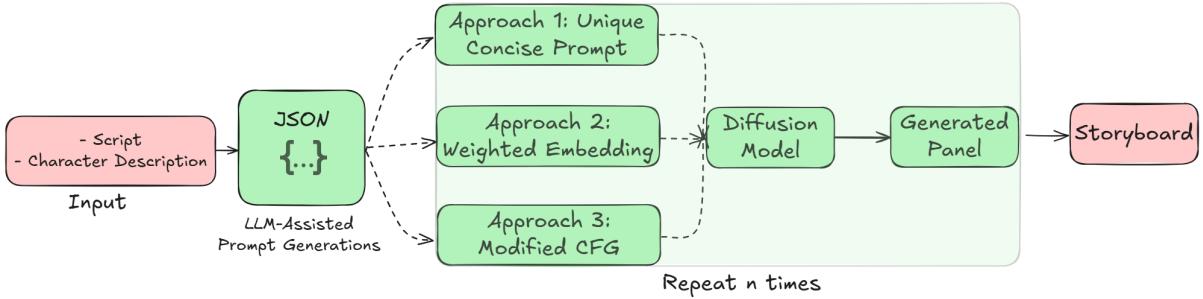


Figure 1: The framework developed to generate the storyboard. From the input, an LLM generates a JSON with all the info necessary for each panel of the storyboard, e.g. shot type, environment, actions, characters, etc. Then, we choose a generation approach to process the prompts coming from the LLM and create the panel. This is iterated for all the n boards that need to be created to get the storyboard.

main advancement lies in the introduction of latent diffusion models, which operate in a compressed latent space rather than pixel space. This approach reduces computational requirements while maintaining high-fidelity output. The model architecture employs a variational autoencoder (VAE) for dimensionality reduction, followed by a U-Net that performs iterative denoising steps conditioned on text embeddings from a pre-trained text encoder.

- **Classifier-Free Guidance (CFG):** Classifier-free guidance [2] is a technique used to enhance text-image alignment in diffusion models by **blending conditional and unconditional predictions**. In the standard CFG framework, the noise prediction at each diffusion timestep is computed as:

$$\hat{\epsilon} = \hat{\epsilon}_{\text{uncond}} + s (\hat{\epsilon}_{\text{cond}} - \hat{\epsilon}_{\text{uncond}}), \quad (1)$$

where $\hat{\epsilon}_{\text{uncond}}$ is the noise prediction without any conditioning (i.e., unconditional), $\hat{\epsilon}_{\text{cond}}$ is the noise prediction conditioned on the text prompt, and s is the guidance scale that controls the strength of conditioning. A higher s value forces the generated image to adhere more closely to the text prompt, at the potential cost of reduced diversity, while a lower s value allows for more creative variations.

3 Methodology

3.1 Input Data and Hierarchical Story Decomposition

Our pipeline takes two inputs: (1) a **script** S , (2) a **character dictionary** C with names and physical attributes of each of them, e.g. age, gender, ethnicity, etc.

We use Meta’s LLM, **Llama 3.1 8B Instruct Turbo** [8] to decompose the input script into n storyboard panels, where n is chosen by the LLM itself. For each panel P_i ($i \in [n]$), the LLM generates the following structured elements:

- **Camera Orientation:** Specification of viewpoint (e.g., front view, three-quarter profile, overhead angle).
- **Shot Type:** Technical framing description (e.g., extreme close-up, medium shot, establishing shot).
- **Character Enumeration:** List of characters present in the scene.
- **Environmental Context:** Detailed description of the setting, including architectural elements, lighting conditions, and atmospheric qualities.
- **Action Description:** Precise depiction of character movements, gestures, expressions, and interactions.

Formally, we can define the extraction process as $\mathcal{E}_{\text{LLM}} : \mathcal{S} \times \mathcal{C} \rightarrow \{P_1, P_2, \dots, P_n\}$. In practice, the output is formatted as a JSON, such as:

```
{
  "scenes": [
    {
      "scene_number": 1,
      "shot_type": "Medium Shot",
      "orientation": "Front View",
      "characters": [
        {"name": "Don Vito Corleone"},
        {"name": "Johnny Fontane"},
        {"name": "Tom Hagen"}
      ],
      "environment": "Don's office, daytime, ...",
      "description": "Don Corleone stands..."
    },
    {
      "scene_number": 2,
      ...
    }
    ...
    {
      "scene_number": n,
      ...
    }
  ]
}
```

Thus, for each scene we end up with the following prompts: s_1 for the **shot type**, s_2 for the **orientation**, $s_{3,i}$ for each **character** i (combined with the description dictionary), s_4 for the **environment**, s_5 for the **description**.

We also define an artistic **style** prompt typical of storyboards s_6 —"rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style"— and a **negative** prompt "low quality, photorealistic, 3d render, overly detailed, digital art, painting, vibrant colors, fine art, NSFW" to steer the synthesis towards the desired rough sketch aesthetic.

3.2 Image Generation Approaches

We employ Stable Diffusion v1.5 as the architecture for image synthesis. We use it because it's small and can run locally. To overcome conventional token length constraints (77 tokens) and effectively emphasize multiple narrative aspects, we investigate three methodological approaches:

Approach 1: Unique Concise Prompting

In this baseline method, a single prompt is constructed by adding the elements of a scene in the following order:

$$\text{Prompt} = s_6 \oplus s_1 \oplus s_2 + \bar{s}_5, \quad (2)$$

where \oplus denotes string concatenation. Notice that to ensure being within the 77 token limit with high probability, just for this approach we use a shortened version of the style prompt, thus having \bar{s}_5 —"rough b&w simple pencil sketch, J.C. Leyendecker style".

This approach directly feeds the composed prompt to the diffusion model, but by construction omits character details and environment descriptions.

Approach 2: Weighted Token Embeddings

To include all extracted information, in this method, each prompt s_i is independently encoded into its own text embedding vector \mathbf{E}_i . A composite embedding for each P_i is then computed via a weighted average:

$$\mathbf{E}_{\text{combined}} = \frac{\sum_i w_i \mathbf{E}_i}{\sum_i w_i}, \quad (3)$$

where $w_i \in \mathbb{R}^+$ are pre-determined weights that modulate the influence of each component. Then $\mathbf{E}_{\text{combined}}$ is fed to the diffusion model for the panel generation. This formulation allows for finer control over the narrative emphasis during image generation, ensuring that even less dominant elements can be amplified or attenuated as needed.

Approach 3: Multi-Prompt Classifier-Free Guidance (CFG)

Extending the standard classifier-free guidance framework, this approach computes independent conditional noise predictions for each semantic component. Denoting the prediction for the i -th component as $\hat{\epsilon}_{\text{cond}_i}$ with associated weight w_i , these predictions are aggregated into a combined conditional term:

$$\hat{\epsilon}_{\text{cond_combined}} = \frac{\sum_i w_i \hat{\epsilon}_{\text{cond}_i}}{\sum_i w_i}. \quad (4)$$

The final noise prediction is then given by:

$$\hat{\epsilon} = \hat{\epsilon}_{\text{uncond}} + s (\hat{\epsilon}_{\text{cond_combined}} - \hat{\epsilon}_{\text{uncond}}), \quad (5)$$

where $\hat{\epsilon}_{\text{uncond}}$ is the unconditional noise prediction, s is the guidance scale parameter, and the subsequent denoising step is performed as:

$$\mathbf{z}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{z}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \hat{\epsilon} \right) + \sigma_t \boldsymbol{\eta}, \quad (6)$$

with α_t and $\bar{\alpha}_t$ as diffusion schedule parameters, σ_t as the noise scale at timestep t , and $\boldsymbol{\eta} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. This approach enables precise control over each narrative component’s contribution while maintaining overall stylistic coherence.

4 Additional Methods

We also explored several alternative methodologies to further enhance the generation process. Although these approaches were not developed in detail, we briefly summarize them for completeness.

4.1 Classifier Guidance for Shot-Type Specification

To ensure that the generated images closely adhere to the specified shot type, we experimented with an auxiliary shot-type classifier. Specifically, we fine-tuned an EfficientNet B3 [7] on a publicly available shot-type dataset [3]. This classifier, denoted by $C_\phi(\mathbf{z}, \text{shot_type})$, is integrated into the diffusion process to modify the noise prediction as follows:

$$\hat{\epsilon}_{\text{guided}} = \hat{\epsilon} + s \nabla_{\mathbf{z}} \log C_\phi(\mathbf{z}, \text{shot_type}), \quad (7)$$

where s controls the guidance strength. Although the classifier reached an accuracy of 75%, the qualitative improvements were unsatisfactory even for simple prompts (see Figure 2). Thus, we concluded that the inherent shot-type capturing capability of Stable Diffusion, despite occasional imperfections, was sufficient for our task, and the additional classifier guidance was unnecessary.

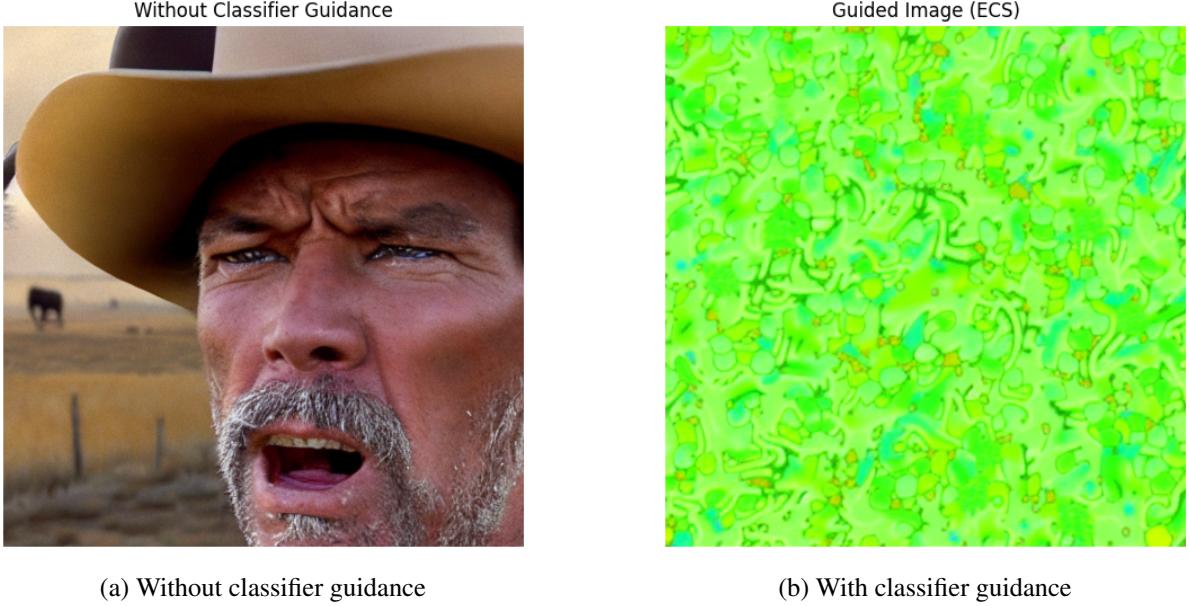


Figure 2: Comparison of generated images for the prompt: “Close up, angry cowboy walks toward his wooden ranch in the Wild West.”

4.2 Character Sheet Generation via ControlNet and DreamBooth

To achieve consistent character appearances across storyboard panels, we implemented a two-stage approach:

1. **Reference Generation:** We used ControlNet [10] conditioned on pose templates to generate reference character illustrations.
2. **Model Fine-Tuning:** We then fine-tuned a DreamBooth model [5] on these reference images, incorporating a unique identifier token into the generation prompts.

While the reference generation step was straightforward, we did not pursue multi-token fine-tuning for DreamBooth due to time constraints. Nonetheless, we believe that further refinement in this direction could significantly improve character consistency across panels.

4.3 Multiple Unconditional Passes

This alternative approach computes individual unconditional passes for each prompt component. Formally, given m subprompts, the final noise prediction is computed as:

$$\hat{\epsilon} = \hat{\epsilon}_{\text{uncond}} + s \sum_{i=1}^m w_i (\hat{\epsilon}_{\text{cond}_i} - \hat{\epsilon}_{\text{uncond}_i}). \quad (8)$$

Although this formulation provides finer control over each prompt component compared to Approach 3, it has two major drawbacks:

1. **Computational Overhead:** Each diffusion step requires $1 + 2m$ UNet calls (one global unconditional pass and two calls per subprompt), compared to only $1 + m$ calls in Approach 3, where each subprompt shared the same baseline unconditional pass. This results in a significantly slower generation.
2. **Model Limitations:** The method relies on generating an unconditional subprompt for each conditional subprompt.

Due to these limitations, and upon seeing that the images did not match the prompts at all, we opted not to pursue further development of this approach.

5 Experiments

5.1 Setup

Our experiments utilize script excerpts from four different movies with associated character dictionaries, as described in Section 3. The following implementation details summarize the experimental setup:

- **Model Architecture:** Stable Diffusion v1.5 (loaded from `runwayml`), with a CLIP ViT-L/14 text encoder.
- **Scheduler:** UniPC Multistep Scheduler[11], with 50 inference steps.
- **Resolution:** All generated panels have a resolution of 512×512 pixels.
- **Guidance Scale:** During CFG we use a standard guidance scale of $s = 7.5$.
- **Prompt Weights:** For Approach 2 and 3, the weights w_i are empirically set as follows:
 - **Style:** 2.0
 - **Shot (Cinematography):** 1.2
 - **Characters:** 0.9 (same for all characters in the panel)
 - **Environment:** 1.0
 - **Description:** 1.5

5.2 Results

Our evaluation is purely qualitative. First of all, we consider how well the LLM can extract scene information that match the script. Doing manual inspections, our model is well capable of following the prompt and define consistent cinematic sequences, which would be indeed done in real life too. Although the model clearly has no cinema knowledge, the stylistic continuity of the shot types and action descriptions are good enough for a storyboard.

As for the image generations, we consider several aspects such as narrative coherence, character consistency, stylistic adherence, and compositional accuracy. **We attach the images and extracted prompts for each script in the Appendix.**

Approach 1: Unique Prompting

This approach consistently yields usable results across all movies. In most cases, it provides stylistically coherent outputs that match the intended sketch-like aesthetic and maintain decent compositional framing.

For instance, in *Titanic* (Figure 4a), this is the only approach that generates reasonable visuals. The panels preserve the sketch style and clearly depict character faces and poses, though close-ups are sometimes overly emphasized and the story is misunderstood. In *The Godfather* (Figure 3a), the generated visuals mostly align with the prompts. Some consistency issues persist—characters lack defining features and the backgrounds are often generic—but narrative flow is preserved, and panel compositions follow logical story progression. Importantly, the sketch style is generally respected. *The Dictator* (Figure 5a) benefits most from this method. The visuals are clean, humorous, and resemble handcrafted political cartoons, aligning well with the film’s satirical tone. The characters are stylized but expressive, and the approach gives a relatively strong sense of story progression. Still, the consistency across images is not good. *Oppenheimer* (Figure 6a) also fares well. Despite occasional issues with overly clean rendering (not always sketch-like), the method captures both environmental context and dramatic intensity, especially in panels portraying the Trinity test and Oppenheimer’s internal tension.

In summary, this approach works well as a default. However, due to token constraints, it sacrifices fine-grained details like character descriptions or environmental subtleties.

Approach 2: Weighted Embeddings

Despite the theoretical flexibility of this method, results show a decline in overall quality. The assumption that embedding blending improves semantic richness does not consistently hold in practice.

In *Titanic* (Figure 4b), outputs are particularly weak. Some panels are even black as they were detected as NSFW content, and many panels show entirely abstract textures—indicating that the embedded prompts may have canceled each other out or diluted important narrative cues. *The Godfather* (Figure 3b) also suffers in this mode. Images tend to revert to generic portraits or stiff full-body shots, often in photographic or realistic styles despite explicit sketch prompts. Character consistency is poor, with different facial features and body types across panels. Close-ups dominate, even in scenes that call for more dynamic framing. *The Dictator* (Figure 5b) performs slightly better, but still inconsistently. Some panels resemble political posters or newspaper-style collages, which are stylistically interesting, but often fail to clearly illustrate the scene’s actions or tone. The lack of control over the weighting may be responsible for these mismatches. In contrast, *Oppenheimer* (Figure 6b) shows relatively better results under Approach 2. Some characters (e.g., Groves or Oppenheimer) recur across panels, and the narrative flow is marginally preserved. However, stylistic adherence fluctuates: several outputs resemble photographs rather than sketches, and the environment is often oversimplified.

We hypothesize that the embedding summation technique may not preserve prompt boundaries clearly, leading to blended representations that confuse the diffusion model. Further tuning of weights or alternative fusion strategies may be necessary for each scene.

Approach 3: Multi-Prompt Classifier-Free Guidance

This approach, developed to improve upon Approach 2, delivers much more structured and faithful storyboards. To make the comparison fair, we decided to keep the same weights as the previous approach.

For *Oppenheimer* (Figure 6c), the sequence is able to convey the same type of narrative arc, even though images lack hand-drawn aesthetics. Some panels did not converge to a meaningful image, probably because of the weight choice. Similarly, *The Dictator* (Figure 5c) yields visually richer results with this method, with consistent costumes and staging. While the sketch style is less evident, panels exhibit a coherent visual story. The model appears to better interpret prompt semantics independently, likely due to explicit guidance separation. In *The Godfather* (Figure 3c), while character fidelity is still limited, the overall shot variety improves over Approach 2, but a better weights adjustment could lead to improved results. *Titanic* (Figure 4c) underperforms here. All panels are abstract and it suggests that, the diffusion model struggled to disentangle overlapping prompt components when guided separately. Alternatively, the separate noise predictions might lack sufficient reinforcement across components (e.g., characters, environment) for such minimal visual scenes.

Overall, Approach 3 works better in structuring and balancing multiple prompt dimensions, but it may struggle stylistically or in low-information visual settings. The quality is really scene and weight-dependent.

6 Discussion and Method of Improvements

The results obtained from our experiments underscore both the potential and the current limitations of diffusion-based approaches for storyboard generation. Overall, Approach 1 demonstrated strong stylistic coherence and narrative flow, while the more complex weighted embeddings and multi-prompt CFG methods exhibited mixed performance with several issues. The additional experiments, including classifier guidance and multiple unconditional passes, further revealed that increased complexity does not automatically translate into improved quality. In many cases, these alternative methods introduced undesirable artifacts or excessive computational overhead without yielding substantial benefits.

These observations suggest several avenues for refinement. First, fine-tuning the diffusion model on a dedicated storyboard or sketch-style dataset may help to naturally enforce the desired aesthetic without the need for extensive auxiliary prompts. Similarly, incorporating architectural insights from recent work

such as [6] could enhance narrative extraction and prompt generation, streamlining the overall process. Another simple improvement lies in using a more advanced image generation model.

A promising direction lies in expanding our use of the DreamBooth framework to achieve better character consistency. By further developing this approach, future work could ensure that distinctive character features are preserved across panels. Moreover, introducing quantitative metrics—such as computing CLIP similarity scores for character segmentation or adapting FID scores to measure distributional similarity with reference storyboards—would offer objective means of evaluating and guiding improvements.

Additional control over storyboard elements is also desirable. Providing users, such as movie writers or directors, with an interactive interface to adjust shot types, camera orientations, and other visual parameters could greatly enhance the practical utility of the system. Finally, extending the framework to generate short video sequences that depict camera movements, potentially with models like AKIRA [9], could further align the output with industry-standard pre-visualization practices.

7 Conclusions

In this work, we presented a novel framework for automated storyboard generation that leverages diffusion models and hierarchical narrative decomposition. By integrating large language models for prompt extraction with a text-to-image synthesis model, our approach is able to generate storyboard panels from movie scripts. Our comparative analysis of three image generation methods revealed that while a baseline approach can capture the desired sketch-like style and narrative flow, more complex methods face challenges in preserving fine details such as character consistency and environmental context, as also demonstrated by our additional experiments in Section 4. These findings underscore the need for further research in several key areas. Future work should explore dedicated fine-tuning on storyboard or sketch-style datasets, expand the DreamBooth-based character consistency approach, and integrate quantitative metrics (e.g., CLIP similarity or adapted FID scores) for objective evaluation. Enhancing user control—allowing interactive adjustments of shot types, camera angles, and other visual parameters—and extending the framework to generate short video sequences depicting camera movements are also promising directions.

Overall, our work lays the foundation for automated storyboard generation and provides valuable insights into the trade-offs and opportunities in diffusion-based visual narrative synthesis.

References

- [1] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks, 2014.
- [2] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance, 2022.
- [3] Anyi Rao, Jiaze Wang, Lining Xu, Xuekun Jiang, Qingqiu Huang, Bolei Zhou, and Dahu Lin. A unified framework for shot type classification based on subject centric lens, 2020.
- [4] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10684–10695, June 2022.
- [5] Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation. 2022.
- [6] Xiaoqian Shen and Mohamed Elhoseiny. Storygpt-v: Large language models as consistent story visualizers, 2023.
- [7] Mingxing Tan and Quoc V. Le. Efficientnet: Rethinking model scaling for convolutional neural networks, 2020.
- [8] Llama Team. The llama 3 herd of models, 2024.
- [9] Xi Wang, Robin Courant, Marc Christie, and Vicky Kalogeiton. Akira: Augmentation kit on rays for optical video generation, 2024.
- [10] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models, 2023.
- [11] Wenliang Zhao, Lujia Bai, Yongming Rao, Jie Zhou, and Jiwen Lu. Unipc: A unified predictor-corrector framework for fast sampling of diffusion models, 2023.

Appendix

A The Godfather

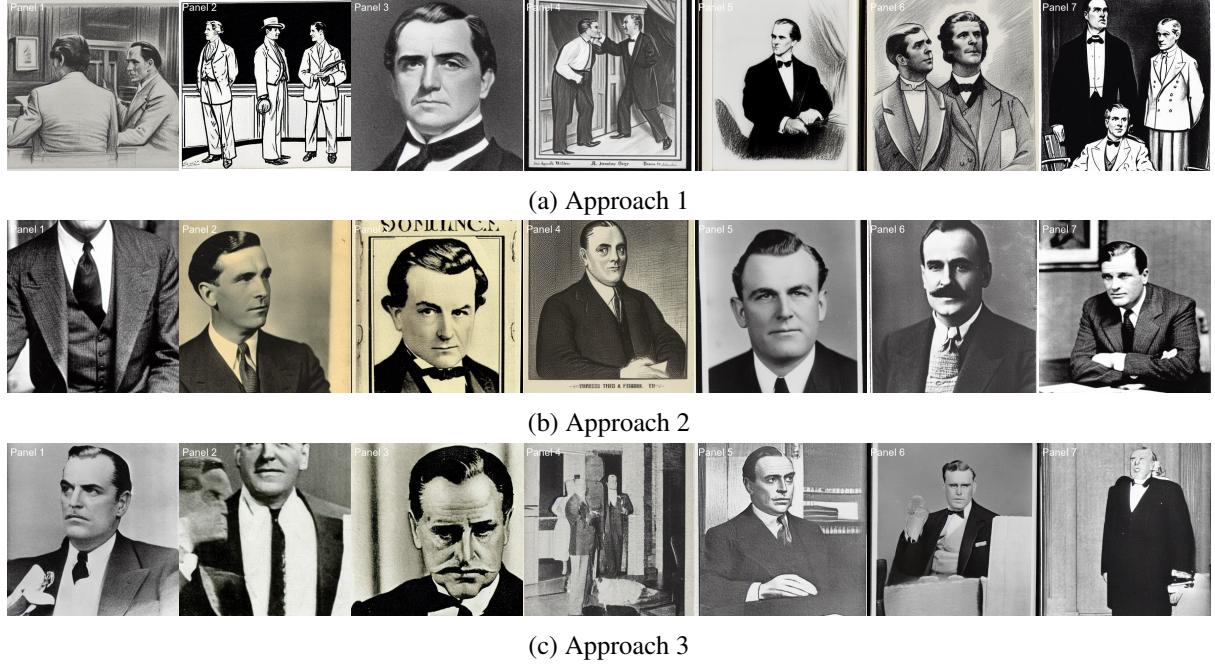


Figure 3: Storyboard generated from a scene in "The Godfather".

A.1 Prompts

Scene 1 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.
- **Subprompt 2 (weight 0.9):** Johnny Fontane: late 30s, male, short, slicked-back black hair, clean shaven, slim and fit, wearing dark, stylish suit with an open collar, with gold ring, cigarette.
- **Subprompt 3 (weight 0.9):** Tom Hagen: German-Irish, early 40s, male, short, neatly combed brown hair, clean-shaven, medium build, upright posture, wearing gray suit, dark tie.
- **Subprompt 4 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 5 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.
- **Subprompt 6 (weight 1.2):** Medium Shot, Front View.
- **Subprompt 7 (weight 1.5):** Don Corleone stands imposingly behind desk, face stern with righteous anger, pointing finger at Johnny. Johnny appears embarrassed, head slightly bowed. Hagen stands to the right, barely containing laughter. Tension and amusement mix in intimate office atmosphere.

Scene 2 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.
- **Subprompt 2 (weight 0.9):** Johnny Fontane: late 30s, male, short, slicked-back black hair, clean shaven, slim and fit, wearing dark, stylish suit with an open collar, with gold ring, cigarette.
- **Subprompt 3 (weight 0.9):** Tom Hagen: German-Irish, early 40s, male, short, neatly combed brown hair, clean-shaven, medium build, upright posture, wearing gray suit, dark tie.
- **Subprompt 4 (weight 0.9):** Sonny: Italian-American, early 30s, male, curly, dark brown hair, clean-shaven, athletic build, wearing formal suit, slightly disheveled.
- **Subprompt 5 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 6 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.
- **Subprompt 7 (weight 1.2):** Medium Shot, Front View.
- **Subprompt 8 (weight 1.5):** Sonny quietly enters room from right, adjusting disheveled clothes. Don leans forward at desk, expression softening to business-like focus. Johnny stands center, straightening posture. Hagen observes from left corner. Atmosphere shifts from personal rebuke to business discussion.

Scene 3 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 3 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.
- **Subprompt 4 (weight 1.2):** Close-Up, Front View.
- **Subprompt 5 (weight 1.5):** Don Corleone's face fills frame, stern and contemplative. Eyes narrowed, jaw set firmly. Saying "I'll make him an offer he can't refuse" with quiet, confident menace. Power and authority emanate from his expression.

Scene 4 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.
- **Subprompt 2 (weight 0.9):** Johnny Fontane: late 30s, male, short, slicked-back black hair, clean shaven, slim and fit, wearing dark, stylish suit with an open collar, with gold ring, cigarette.
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 4 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.

- **Subprompt 5 (weight 1.2):** Medium Close-Up, Three-Quarters View.
- **Subprompt 6 (weight 1.5):** Don Corleone escorts Johnny to door, pinching his cheek firmly. Don's expression shows affection mixed with dominance. Johnny winces slightly at pain while showing relief and gratitude. Door frame visible on right edge of shot.

Scene 5 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.
- **Subprompt 2 (weight 0.9):** Tom Hagen: German-Irish, early 40s, male, short, neatly combed brown hair, clean-shaven, medium build, upright posture, wearing gray suit, dark tie.
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 4 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.
- **Subprompt 5 (weight 1.2):** Medium Shot, Front View.
- **Subprompt 6 (weight 1.5):** Don Corleone turns from closed door, small smile fading to serious business expression. Hagen stands attentively near desk, notepad ready. Don moves toward chair, shoulders slightly hunched, gold ring catching light as he gestures.

Scene 6 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.
- **Subprompt 2 (weight 0.9):** Tom Hagen: German-Irish, early 40s, male, short, neatly combed brown hair, clean-shaven, medium build, upright posture, wearing gray suit, dark tie.
- **Subprompt 3 (weight 0.9):** Sonny: Italian-American, early 30s, male, curly, dark brown hair, clean-shaven, athletic build, wearing formal suit, slightly disheveled.
- **Subprompt 4 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 5 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.
- **Subprompt 6 (weight 1.2):** Over-the-Shoulder Shot, Profile View.
- **Subprompt 7 (weight 1.5):** Camera over Don's shoulder, facing Hagen and Sonny. Don's gray-black hair and dark suit visible in foreground. Hagen's face shows respectful attention. Sonny stands beside him, now composed. Don's voice carries weight as he issues final instructions about hospital visit.

Scene 7 Subprompts:

- **Subprompt 1 (weight 0.9):** Don Vito Corleone: Italian-American, early 60s, male, slicked-back gray-black hair, stocky, slightly hunched posture, wearing dark three-piece suit, with gold ring on right hand, pocket watch.

- **Subprompt 2 (weight 0.9):** Tom Hagen: German-Irish, early 40s, male, short, neatly combed brown hair, clean-shaven, medium build, upright posture, wearing gray suit, dark tie.
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style.
- **Subprompt 4 (weight 1.0):** Don's office, daytime, summer 1945. Elegant wood-paneled room with large desk, leather chairs, warm lighting filtering through venetian blinds.
- **Subprompt 5 (weight 1.2):** Medium Shot, Front View.
- **Subprompt 6 (weight 1.5):** Don Corleone stands near chair, eyes filled with concern. Hagen stands attentively near desk, expression reflecting understanding. Don's hands clasped together as he shares news about Genco Abbandando's hospitalization.

B Titanic

B.1 Prompts

Scene 1 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Point-of-View Shot, Profile View
- **Subprompt 5 (weight 1.5):** Rose stares upwards, hands locked with Jack's, reflecting on stars above.

Scene 2 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Rose's pale face and eyes closed as she sings softly.

Scene 3 Subprompts:

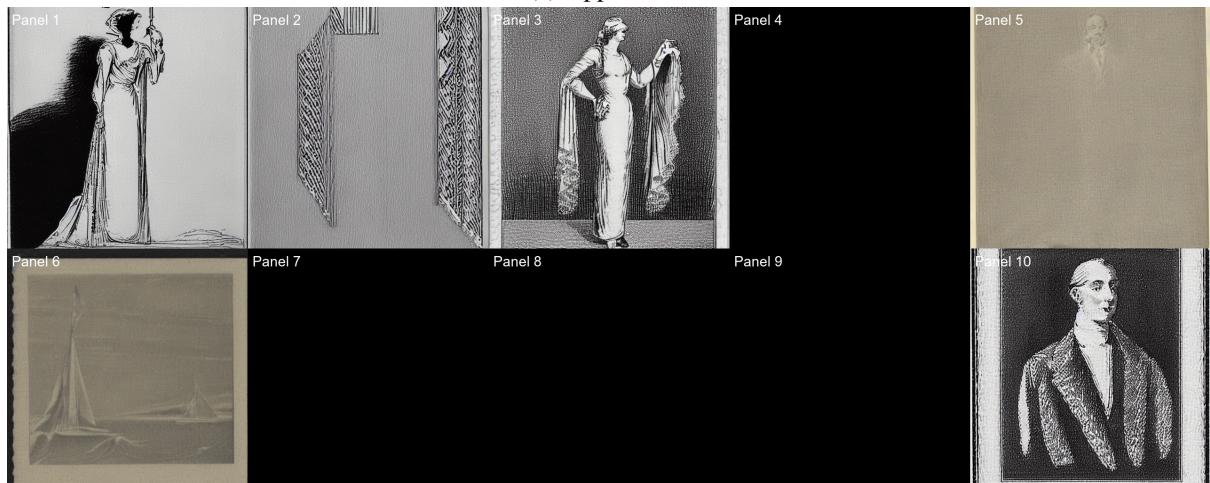
- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Point-of-View Shot, Profile View
- **Subprompt 5 (weight 1.5):** Rose sees the stars in close-up, with Milky Way visible from horizon to horizon.

Scene 4 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style



(a) Approach 1



(b) Approach 2



(c) Approach 3

Figure 4: Storyboard generated from a scene in the "Titanic".

- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Medium Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Rose's face contorted in pain, her breathing shallow.

Scene 5 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 0.9):** Jack: American, early 20s, male, medium length, blonde, clean shaven, fit, wearing shirt, all wet
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 5 (weight 1.2):** Two-Shot, Front View
- **Subprompt 6 (weight 1.5):** Rose reaches out to Jack, his face frozen in frost, her hand locked to his.

Scene 6 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Point-of-View Shot, Profile View
- **Subprompt 5 (weight 1.5):** Rose sees the distant boat in slow motion, its oars lifting out of the water, voices of men sounding slow and distorted.

Scene 7 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Medium Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Rose's face, weak and hopeless, looking at Jack's frozen face.

Scene 8 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Point-of-View Shot, Profile View
- **Subprompt 5 (weight 1.5):** Rose sees the boat in the distance, its torch light a star impossibly far away, she tries to cry out.

Scene 9 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Rose's face filled with determination as she breaks away from Jack's frozen hand.

Scene 10 Subprompts:

- **Subprompt 1 (weight 0.9):** Rose: American, early 20s, female, long, dark brown, fit, wearing dress, all wet
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** EXT. OCEAN 291 - night. Black water with stars reflecting in the surface.
- **Subprompt 4 (weight 1.2):** Medium Long Shot, Front View
- **Subprompt 5 (weight 1.5):** Rose plunges into icy water, grabbing a whistle, blowing it with all her strength.

C The Dictator

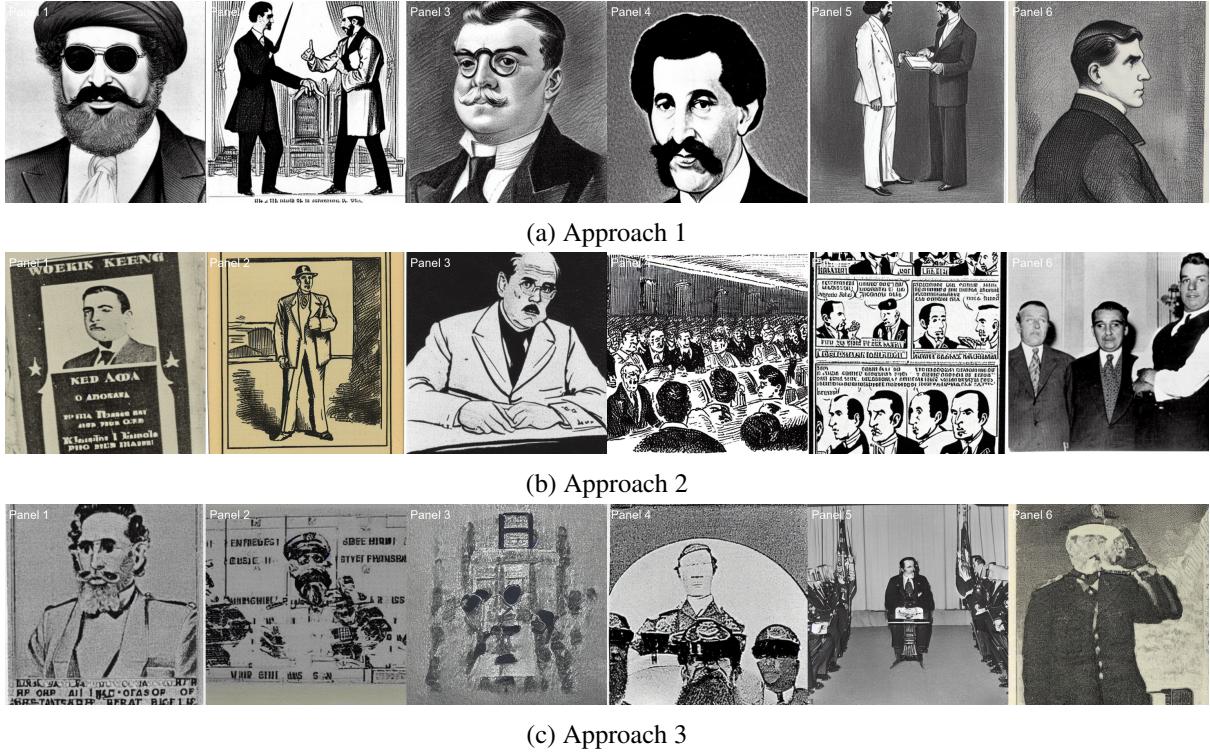


Figure 5: Storyboard generated from a scene in "The Dictator".

C.1 Prompts

Scene 1 Subprompts:

- **Subprompt 1 (weight 0.9):** Supreme Leader Aladeen: late 30s to early 40s, thick, full beard, average build, wearing military uniform with ornate decorations, with sunglasses, military medals
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** Military briefing room, daytime. Dimly lit, with a large screen behind the Supreme Leader and rows of high-tech equipment.
- **Subprompt 4 (weight 1.2):** Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Supreme Leader Aladeen's face fills the frame, thick beard and sunglasses visible. His expression shows a mix of amusement and seriousness.

Scene 2 Subprompts:

- **Subprompt 1 (weight 0.9):** Supreme Leader Aladeen: late 30s to early 40s, thick, full beard, average build, wearing military uniform with ornate decorations, with sunglasses, military medals
- **Subprompt 2 (weight 0.9):** Nuclear Nadal: late 40s, mustache, slim, wearing white lab coat over formal attire, with glasses, clipboard
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style

- **Subprompt 4 (weight 1.0):** Military briefing room, daytime. Dimly lit, with a large screen behind the Supreme Leader and rows of high-tech equipment.
- **Subprompt 5 (weight 1.2):** Two-Shot, Front View
- **Subprompt 6 (weight 1.5):** Supreme Leader Aladeen and Nuclear Nadal stand facing each other, the Supreme Leader holding a pointer. Nadal looks puzzled, clipboard in hand.

Scene 3 Subprompts:

- **Subprompt 1 (weight 0.9):** Nuclear Nadal: late 40s, mustache, slim, wearing white lab coat over formal attire, with glasses, clipboard
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** Military briefing room, daytime. Dimly lit, with a large screen behind the Supreme Leader and rows of high-tech equipment.
- **Subprompt 4 (weight 1.2):** Medium Close-Up, Three-Quarters View
- **Subprompt 5 (weight 1.5):** Nuclear Nadal's face fills the frame, mustache and glasses visible. His expression shows concern and frustration.

Scene 4 Subprompts:

- **Subprompt 1 (weight 0.9):** Supreme Leader Aladeen: late 30s to early 40s, thick, full beard, average build, wearing military uniform with ornate decorations, with sunglasses, military medals
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** Military briefing room, daytime. Dimly lit, with a large screen behind the Supreme Leader and rows of high-tech equipment.
- **Subprompt 4 (weight 1.2):** Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Supreme Leader Aladeen's face fills the frame, a determined look on his face. He says 'kaboom' with a hint of excitement.

Scene 5 Subprompts:

- **Subprompt 1 (weight 0.9):** Supreme Leader Aladeen: late 30s to early 40s, thick, full beard, average build, wearing military uniform with ornate decorations, with sunglasses, military medals
- **Subprompt 2 (weight 0.9):** Nuclear Nadal: late 40s, mustache, slim, wearing white lab coat over formal attire, with glasses, clipboard
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** Military briefing room, daytime. Dimly lit, with a large screen behind the Supreme Leader and rows of high-tech equipment.
- **Subprompt 5 (weight 1.2):** Medium Shot, Front View
- **Subprompt 6 (weight 1.5):** Supreme Leader Aladeen and Nuclear Nadal stand facing each other, the Supreme Leader's expression changing from amusement to seriousness. Nadal looks exasperated, clipboard in hand.

Scene 6 Subprompts:

- **Subprompt 1 (weight 0.9):** Supreme Leader Aladeen: late 30s to early 40s, thick, full beard, average build, wearing military uniform with ornate decorations, with sunglasses, military medals
- **Subprompt 2 (weight 0.9):** Nuclear Nadal: late 40s, mustache, slim, wearing white lab coat over formal attire, with glasses, clipboard
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** Military briefing room, daytime. Dimly lit, with a large screen behind the Supreme Leader and rows of high-tech equipment.
- **Subprompt 5 (weight 1.2):** Over-the-Shoulder Shot, Profile View
- **Subprompt 6 (weight 1.5):** Camera over the Supreme Leader's shoulder, facing Nuclear Nadal. The Supreme Leader's profile and the rows of high-tech equipment visible in the background. Nadal's expression shows frustration and concern.

D Oppenheimer

D.1 Prompts

Scene 1 Subprompts:

- **Subprompt 1 (weight 0.9):** J. Robert Oppenheimer: early 40s, clean-shaven, slim, wearing white shirt, dark trousers, wide-brimmed hat, with pipe
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** Trinity Test Site, desert landscape, evening. Steel tower stands tall in center, atomic bomb at top. Storm clouds gather on horizon.
- **Subprompt 4 (weight 1.2):** Aerial View, Front View
- **Subprompt 5 (weight 1.5):** J. Robert Oppenheimer watches last technician descend from steel tower. Oppenheimer nods at technician, then begins lonely climb up tower. Atomic bomb looms large, adorned with detonators and wires.

Scene 2 Subprompts:

- **Subprompt 1 (weight 0.9):** General Leslie Groves: late 40s, clean-shaven, stocky, wearing military uniform, with military cap
- **Subprompt 2 (weight 0.9):** Weatherman: late 30s, clean-shaven, average build, wearing military uniform, with radio set
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** Base Camp, Trinity Test Site, evening. Army tents scattered about, windmill spinning in distance.
- **Subprompt 5 (weight 1.2):** Long Shot, Front View
- **Subprompt 6 (weight 1.5):** General Leslie Groves and weatherman stand amidst army tents, concern etched on their faces. Groves asks weatherman about possibility of delay due to approaching storm.

Scene 3 Subprompts:

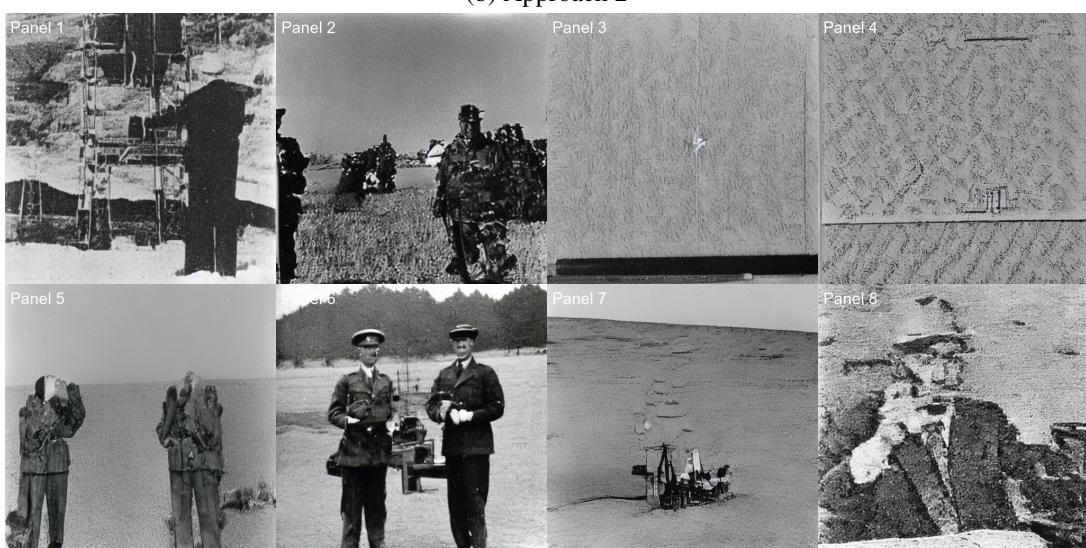
- **Subprompt 1 (weight 0.9):** J. Robert Oppenheimer: early 40s, clean-shaven, slim, wearing white shirt, dark trousers, wide-brimmed hat, with pipe
- **Subprompt 2 (weight 0.9):** George Kistiakowsky: mid-40s, mustache, average build, wearing white lab coat, with clipboard
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** Base Camp, Trinity Test Site, evening. Army tents scattered about, windmill spinning in distance.
- **Subprompt 5 (weight 1.2):** Two-Shot, Front View
- **Subprompt 6 (weight 1.5):** Kistiakowsky bursts into camp, speaking with Oppenheimer. Oppenheimer listens intently, brow furrowed in concern. Kistiakowsky holds out hand, offering bet about implosion lenses working.



(a) Approach 1



(b) Approach 2



(c) Approach 3

Figure 6: Storyboard generated from a scene in "Oppenheimer".

Scene 4 Subprompts:

- **Subprompt 1 (weight 0.9):** Enrico Fermi: mid-40s, clean-shaven, average build, wearing white lab coat over suit, with glasses
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** Base Camp, Trinity Test Site, evening. Army tents scattered about, windmill spinning in distance.
- **Subprompt 4 (weight 1.2):** Close-Up, Front View
- **Subprompt 5 (weight 1.5):** Fermi takes bets from team members about bomb's yield. Fermi's eyes gleam with excitement as he collects wagers, adding to his calculations.

Scene 5 Subprompts:

- **Subprompt 1 (weight 0.9):** J. Robert Oppenheimer: early 40s, clean-shaven, slim, wearing white shirt, dark trousers, wide-brimmed hat, with pipe
- **Subprompt 2 (weight 0.9):** General Leslie Groves: late 40s, clean-shaven, stocky, wearing military uniform, with military cap
- **Subprompt 3 (weight 0.9):** Weatherman: late 30s, clean-shaven, average build, wearing military uniform, with radio set
- **Subprompt 4 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 5 (weight 1.0):** South Observation Post, Trinity Test Site, night. Rain pounds against desert floor, lightning flashes across sky.
- **Subprompt 6 (weight 1.2):** Medium Shot, Front View
- **Subprompt 7 (weight 1.5):** Oppenheimer, Groves, and weatherman stand together, observing storm. Groves asks about possibility of delay, weatherman's expression grim. Oppenheimer offers insight about desert weather patterns.

Scene 6 Subprompts:

- **Subprompt 1 (weight 0.9):** General Leslie Groves: late 40s, clean-shaven, stocky, wearing military uniform, with military cap
- **Subprompt 2 (weight 0.9):** J. Robert Oppenheimer: early 40s, clean-shaven, slim, wearing white shirt, dark trousers, wide-brimmed hat, with pipe
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** South Observation Post, Trinity Test Site, night. Rain pounds against desert floor, lightning flashes across sky.
- **Subprompt 5 (weight 1.2):** Long Shot, Front View
- **Subprompt 6 (weight 1.5):** Groves and Oppenheimer step outside into rain, Frank approaching them. Frank announces arming party's departure, evacuation preparations underway.

Scene 7 Subprompts:

- **Subprompt 1 (weight 0.9):** General Leslie Groves: late 40s, clean-shaven, stocky, wearing military uniform, with military cap
- **Subprompt 2 (weight 0.9):** J. Robert Oppenheimer: early 40s, clean-shaven, slim, wearing white shirt, dark trousers, wide-brimmed hat, with pipe
- **Subprompt 3 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 4 (weight 1.0):** Trinity Test Site, pre-dawn. Desert landscape stretches out, silence broken only by distant thunder.
- **Subprompt 5 (weight 1.2):** Extreme Long Shot, Front View
- **Subprompt 6 (weight 1.5):** Groves and Oppenheimer stand together, waiting for detonation. Countdown begins, tension builds. Atomic bomb looms in distance, Oppenheimer's eyes fixed on it.

Scene 8 Subprompts:

- **Subprompt 1 (weight 0.9):** J. Robert Oppenheimer: early 40s, clean-shaven, slim, wearing white shirt, dark trousers, wide-brimmed hat, with pipe
- **Subprompt 2 (weight 2):** rough b&w pencil sketch, simple sketch lines, minimal shading, rough hatching, draft-style, J.C. Leyendecker style
- **Subprompt 3 (weight 1.0):** Trinity Test Site, post-detonation. Mushroom cloud rises high into sky, desert landscape scarred.
- **Subprompt 4 (weight 1.2):** Aerial View, Front View
- **Subprompt 5 (weight 1.5):** Oppenheimer's face contorted in anguish, eyes wide with horror. "Now I am become Death... the destroyer of worlds." whispered in stunned silence.